

Fundamentals of Data Science

WEEK 7: Mathematical Computation Using NumPy

Prepared by: Krishna Devkota

The NumPy array object

What are NumPy and NumPy arrays?

- High-level number objects: integers and floating point
- Containers: lists (costless insertion and append)
- Numpy provides extension package to Python for multi-dimensional arrays
- Closer to hardware (efficiency)
- Designed for scientific computation (convenience)
- Also known as array oriented computing

```
In [1]: 1 import numpy as np
        2
        3 print(np.__version__)
```

1.20.3

```
In [2]: 1 import numpy as np # recommended import conventions
        2
        3 a = np.array([0,1,2,3])
        4
        5 type(a)
```

Out[2]: numpy.ndarray

Why is NumPy useful?

It is a memory-efficient container that provides fast numerical operations.

Interactive Help

```
In [3]: 1 np.array?
```

```
In [4]: 1 np.lookfor("dot")
```

Search results for 'dot'

numpy.dot

Dot product of two arrays. Specifically,

numpy.vdot

Return the dot product of two vectors.

numpy.tensordot

Compute tensor dot product along specified axes.

numpy.ma.dot

Return the dot product of two arrays.

numpy.chararray.dot

Dot product of two arrays.

numpy.linalg.multi_dot

Compute the dot product of two or more arrays in a single function call,

numpy.random.SFC64

BitGenerator for Chris Doty-Humphrey's Small Fast Chaotic PRNG.

numpy.linalg.tensorinv

Compute the 'inverse' of an N-dimensional array.

numpy.ma.MaskedArray.dot

Method dot: product of two arrays. Note that 'a.dot(b)' and 'numpy.dot(a,b)' are

```
In [5]: 1 np.con*?
```

Import conventions

```
In [6]: 1 import numpy as np
```

Creating arrays

Manual construction of arrays

1-D:

```
In [7]: 1 a = np.array([0, 1, 2, 3])
        2 a
```

```
Out[7]: array([0, 1, 2, 3])
```

```
In [8]: 1 a.ndim
```

```
Out[8]: 1
```

```
In [9]: 1 a.shape
```

```
Out[9]: (4,)
```

```
In [10]: 1 len(a)
```

```
Out[10]: 4
```

2-D, 3-D:

```
In [11]: 1 b = np.array([[0, 1, 2], [3, 4, 5]])
        2 b
```

```
Out[11]: array([[0, 1, 2],
                [3, 4, 5]])
```

```
In [12]: 1 b.ndim
```

```
Out[12]: 2
```

```
In [13]: 1 b.shape
```

```
Out[13]: (2, 3)
```

```
In [14]: 1 len(b)
```

```
Out[14]: 2
```

```
In [15]: 1 c = np.array([[1], [2]], [[3], [4]])
        2 c
```

```
Out[15]: array([[1],
                [2]],
               [[3],
                [4]])
```

```
In [16]: 1 c.shape
```

```
Out[16]: (2, 2, 1)
```

```
In [17]: 1 b = np.array([[1,4], [2,7], [4,3]], [[0, 1], [3, 4],[5,6]] )
        2 b.shape
```

```
Out[17]: (2, 3, 2)
```

Exercise

- Create a simple two dimensional array. First, redo the examples from above. And then create your own: how about odd numbers counting backwards on the first row, and even numbers on the second?
- Use the functions `len()`, `numpy.shape()` on these arrays. How do they relate to each other? And to the `ndim` attribute of the arrays?

Functions for creating arrays

In practice, we rarely enter the items one by one.

Create evenly spaced arrays:

```
In [18]: 1 a = np.arange(10) # 0 .. n-1
          2 print(a)

[0 1 2 3 4 5 6 7 8 9]
```

```
In [19]: 1 b = np.arange(1, 9, 2) # start, end (exclusive), step
          2 print(b)

[1 3 5 7]
```

Create array by specifying number of points:

```
In [20]: 1 c = np.linspace(0, 1, 5) # start, end, num-points
          2 print(c)

[0.  0.25 0.5  0.75 1.  ]
```

```
In [21]: 1 d = np.linspace(0, 1, 5, endpoint=False)
          2 print(d)

[0.  0.2 0.4 0.6 0.8]
```

```
In [22]: 1 # np.linspace?
```

Some common arrays:

Arrays with ones

```
In [23]: 1 a = np.ones((3, 3)) # reminder: (3, 3) is a tuple
          2 print(a)

[[1. 1. 1.]
 [1. 1. 1.]
 [1. 1. 1.]]
```

Arrays with zeros (Zero Matrix)

```
In [24]: 1 b = np.zeros((3, 3))
          2 print(b)

[[0. 0. 0.]
 [0. 0. 0.]
 [0. 0. 0.]]
```

Unit Matrix (Identity Matrix)

```
In [25]: 1 c = np.eye(3)
          2 print(c)

[[1. 0. 0.]
 [0. 1. 0.]
 [0. 0. 1.]]
```

Diagonal Matrix

```
In [26]: 1 d = np.diag(np.array([11, 27, 34, 4]))
          2 print(d)

[[11  0  0  0]
 [ 0 27  0  0]
 [ 0  0 34  0]
 [ 0  0  0  4]]
```

Random numbers

```
In [27]: 1 a = np.random.rand(70)      # uniform in [0, 1]
          2 print(a)

[0.50755507 0.0211933  0.43352176 0.44631306 0.23881999 0.83024573
 0.74476418 0.586479   0.49286785 0.48735588 0.2667407  0.6050111
 0.75354372 0.27058423 0.52230328 0.09832853 0.71363667 0.88404059
 0.56705442 0.99448158 0.17873977 0.01220009 0.45699848 0.93175194
 0.84602469 0.47332988 0.90255503 0.22599553 0.30415374 0.71499388
 0.72409148 0.01867644 0.2858131  0.58048634 0.93078663 0.3389969
 0.12008312 0.51627271 0.69920706 0.29864068 0.86160962 0.9058072
 0.76858325 0.26123164 0.9384556  0.93864246 0.74504455 0.91073504
 0.23722471 0.49496735 0.80987834 0.95456578 0.63748325 0.91084975
 0.69213675 0.04294299 0.8335869  0.36994852 0.936557   0.48305288
 0.12533161 0.96445418 0.01702583 0.67657077 0.14043997 0.15531285
 0.64955775 0.98165422 0.69480742 0.76197389]
```

```
In [28]: 1 b = np.random.randn(70)     # Gaussian
          2 print(b)

[-1.08057716 -0.22380444  0.68662417 -0.73620379 -0.01506303 -0.67925245
  0.46199332 -1.12674566  2.11978867  1.61926528  0.6148901  -0.88718431
 -1.06112917  0.86899734  0.0660828  0.8491745  -1.8879036  0.61746113
 -1.26026979 -1.21270503  0.81937935  1.83003338  0.06174604 -0.64577448
 -0.02463203  0.06691692 -1.03999184 -0.2560261  -0.23238899 -1.00755283
 -0.96900851  0.3094744  -0.11024096  1.60897676 -0.75604294  1.02829799
 -0.02520225 -0.33355825 -0.3729897  0.17476908 -1.19253672  1.30436895
 -1.41566081  0.47923289  0.23361135 -0.12545121 -0.28385152  0.26110561
  0.82626166  0.31635129  0.11070639  1.70075073  1.07674992  0.06667967
  0.01521925 -0.03649094 -1.27846407 -0.09173462 -0.44781851 -0.29782149
 -0.5062295  -0.1050317  -0.50717154  0.50516987  0.85546728  0.18512002
  1.1571755  -1.07224293  1.34049442  1.41773681]
```

Seeding a random number generator

```
In [29]: 1 import numpy as np
          2 np.random.seed(42)          # Setting the random seed
          3
          4 b = np.random.rand(3)
          5 print(b)

[0.37454012 0.95071431 0.73199394]
```

Exercise

- Experiment with `arange`, `linspace`, `ones`, `zeros`, `eye` and `diag`.
- Create different kinds of arrays with random numbers.
- Try setting the seed before creating an array with random values.
- Look at the function `np.empty`. What does it do? When might this be useful?

Basic Data Types

You might sometimes notice that, in some instances, array elements are displayed with a trailing dot (e.g. 2. vs 2). This is due to a difference in the data-type used:

```
In [30]: 1 import numpy as np
2 a = np.array([1, 2, 3])
3 print(a.dtype)
4 print(a)
```

```
int32
[1 2 3]
```

```
In [31]: 1 b = np.array([1., 2., 3.])
2 print(b.dtype)
3 print(b)
```

```
float64
[1. 2. 3.]
```

Different data-types allow us to store data more compactly in memory, but most of the time we simply work with floating point numbers. Note that, in the example above, NumPy auto-detects the data-type from the input.

You can explicitly specify which data-type you want:

```
In [32]: 1 c = np.array([1, 2, 3], dtype=float)
2 print(c.dtype)
3 print(c)
```

```
float64
[1. 2. 3.]
```

The **default** data type is floating point:

```
In [33]: 1 a = np.ones((3, 3))
2 print(a.dtype)
```

```
float64
```

Other Data Types

Complex:

```
In [34]: 1 d = np.array([1+2j, 3+4j, 5+6*1j])
2 print(d.dtype)
3 print(d)
```

```
complex128
[1.+2.j 3.+4.j 5.+6.j]
```

Bool:

```
In [35]: 1 e = np.array([True, False, False, True])
2 print(e.dtype)
3 print(e)
```

```
bool
[ True False False  True]
```

Strings:

```
In [36]: 1 f = np.array(['Bonjour', 'Hello', 'Hallo'])
2 print(f.dtype)# <--- strings containing max. 7 letters
3 print(f)
```

```
<U7
['Bonjour' 'Hello' 'Hallo']
```

Others:

- int32
- int64
- uint32
- uint64


Basic visualization

We will cover visualization in detail in the coming session, but let us look at some of the basics below:

```
In [37]: 1 %matplotlib inline
          2 import matplotlib.pyplot as plt
```


```
In [38]: 1 x = np.linspace(-3, 3, 20)
          2 y = np.linspace(0, 9, 20)
          3 plt.plot(x, y**2)      # Line plot
```

Out[38]: [matplotlib.lines.Line2D at 0x1e4466c3a00>]



```
In [39]: 1 plt.plot(x, y, '.')    # dot plot
```

Out[39]: [matplotlib.lines.Line2D at 0x1e446e581c0>]



Plotting 2D arrays


```
In [40]: 1 import numpy as np
          2 import matplotlib.pyplot as plt
```

```
In [41]: 1 image = np.random.rand(1)
          2 print(image)
```

[0.59865848]


```
In [42]: 1 image = np.random.rand(3,3)
2 print(image.shape)
3 print(image)
4 plt.imshow(image, cmap=plt.cm.gray)
5 plt.colorbar();
```

```
(3, 3)
[[0.15601864 0.15599452 0.05808361]
 [0.86617615 0.60111501 0.70807258]
 [0.02058449 0.96990985 0.83244264]]
```



```
In [43]: 1 # Extra
2 a = np.linspace(-5,5,100)
3 plt.plot(a, np.cos(a))
4 plt.plot(a, np.sin(a))
```

Out[43]: [



Exercise: Simple visualizations

- Plot some simple arrays: a cosine as a function of time and a 2D matrix.
- Try using the `copper` colormap on the 2D matrix.

Indexing and slicing

The items of an array can be accessed and assigned to the same way as other Python sequences (e.g. lists):

```
In [44]: 1 a = np.arange(10)
2 print(a)
```

```
[0 1 2 3 4 5 6 7 8 9]
```

```
In [45]: 1 a[0], a[2], a[-1]
```

Out[45]: (0, 2, 9)

The usual python idiom for reversing a sequence is supported:

```
In [46]: 1 a[::-1]
```

Out[46]: array([9, 8, 7, 6, 5, 4, 3, 2, 1, 0])

For multidimensional arrays, indexes are tuples of integers:

```
In [47]: 1 a = np.diag([7,9,6])
         2 a
```

```
Out[47]: array([[7, 0, 0],
               [0, 9, 0],
               [0, 0, 6]])
```

```
In [48]: 1 a[1,1]
```

```
Out[48]: 9
```

```
In [49]: 1 a[2, 1] = 10 # third line, second column
         2 a
```

```
Out[49]: array([[ 7,  0,  0],
               [ 0,  9,  0],
               [ 0, 10,  6]])
```

```
In [50]: 1 a[1]
```

```
Out[50]: array([0, 9, 0])
```

Note:

- In 2D, the first dimension corresponds to rows, the second to columns.
- for multidimensional a, `a[0]` is interpreted by taking all elements in the unspecified dimensions.

Slicing: Arrays, like other Python sequences can also be sliced:

```
In [51]: 1 a = np.arange(10)
         2 a
```

```
Out[51]: array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])
```

```
In [52]: 1 a[2:7:2] # [start:end:step]
```

```
Out[52]: array([2, 4, 6])
```

```
In [53]: 1 a[:4]
```

```
Out[53]: array([0, 1, 2, 3])
```

```
In [54]: 1 a[1:3]
         2 a[::2]
         3 # a[3:]
```

```
Out[54]: array([0, 2, 4, 6, 8])
```

An illustrated summary of NumPy indexing and slicing:

```
>>> a[0,3:5]
array([3,4])
```

```
>>> a[4:,4:]
array([[44, 45],
       [54, 55]])
```

```
>>> a[:,2]
array([2,12,22,32,42,52])
```

```
>>> a[2::2,::2]
array([[20,22,24]
       [40,42,44]])
```

| | | | | | |
|----|----|----|----|----|----|
| 0 | 1 | 2 | 3 | 4 | 5 |
| 10 | 11 | 12 | 13 | 14 | 15 |
| 20 | 21 | 22 | 23 | 24 | 25 |
| 30 | 31 | 32 | 33 | 34 | 35 |
| 40 | 41 | 42 | 43 | 44 | 45 |
| 50 | 51 | 52 | 53 | 54 | 55 |

You can also combine assignment and slicing:

```
In [55]: 1 a = np.arange(10)
          2 print(a)
          3 a[5:] = 7
          4 print(a)
```

```
[0 1 2 3 4 5 6 7 8 9]
[0 1 2 3 4 7 7 7 7 7]
```

```
In [56]: 1 b = np.arange(5)
          2 print(b)
          3 a[5:] = b[::-1]
          4 print(a)
```

```
[0 1 2 3 4]
[0 1 2 3 4 4 3 2 1 0]
```

Exercise: Indexing and Slicing

- Try the different flavours of slicing, using `start`, `end` and `step`: starting from a `linspace`, try to obtain odd numbers counting backwards, and even numbers counting forwards.

Copies and views

A slicing operation creates a view on the original array, which is just a way of accessing array data. Thus the original array is not copied in memory. You can use `np.may_share_memory()` to check if two arrays share the same memory block. Note however, that this uses heuristics and may give you false positives.

```
In [57]: 1 a = np.arange(10)
          2 print(a)
```

```
[0 1 2 3 4 5 6 7 8 9]
```

```
In [58]: 1 b = a[::2]
          2 print(b)
```

```
[0 2 4 6 8]
```

```
In [59]: 1 np.may_share_memory(a, b)
```

Out[59]: True

```
In [60]: 1 b[0] = 12
          2 print(b)
```

```
[12 2 4 6 8]
```

```
In [61]: 1 print(a)
```

```
[12 1 2 3 4 5 6 7 8 9]
```

```
In [62]: 1 a = np.arange(10)
          2 c = a[::2].copy() # force a copy
          3 c[0] = 12
          4 print(a)
```

```
[0 1 2 3 4 5 6 7 8 9]
```

```
In [63]: 1 np.may_share_memory(a, c)
```

Out[63]: False

This behavior can be surprising at first sight... but it allows to save both memory and time.

Numerical operations on arrays

Elementwise operations

Basic operations

With scalars:

```
In [64]: 1 a = np.array([1, 2, 3, 4])
2
3 print(a + 1)
4 print(2**a)
```

```
[2 3 4 5]
[ 2  4  8 16]
```

All arithmetic operates elementwise:

```
In [65]: 1 b = np.ones(4) + 1
2 print(a - b)
3
4 j = np.arange(5)
5 print(2**(j + 1) - j)
```

```
[-1.  0.  1.  2.]
[ 2  3  6 13 28]
```

These operations are of course much faster than if you did them in pure python:

```
In [66]: 1 a = np.arange(10000)
2 %timeit a + 1
3
4 l = range(10000)
5 %timeit [i+1 for i in l]
```

```
7.03 µs ± 306 ns per loop (mean ± std. dev. of 7 runs, 100000 loops each)
694 µs ± 47.8 µs per loop (mean ± std. dev. of 7 runs, 1000 loops each)
```

Array multiplication is not matrix multiplication:

```
In [67]: 1 c = np.ones((3, 3))
2
3 print(c * c) # NOT matrix multiplication!
```

```
[[1. 1. 1.]
 [1. 1. 1.]
 [1. 1. 1.]]
```

Note: Matrix multiplication:

```
In [68]: 1 print(c.dot(c))
```

```
[[3. 3. 3.]
 [3. 3. 3.]
 [3. 3. 3.]]
```

A function to compute the dotproduct of two given arrays

```
In [69]: 1 def dotproduct(a,b):
2     if a.shape[1]!=b.shape[0]:
3         print("Array multiplication not possible for given dimensions")
4     else:
5         return (a.dot(b))
6
7 a=np.ones((1,2))
8 b=np.ones((2,3))
9
10 print(a,b)
11 print(dotproduct(a,b))
```

```
[[1. 1.]] [[1. 1. 1.]
 [1. 1. 1.]]
[[2. 2. 2.]]
```

Exercise: Elementwise operations

- Try simple arithmetic elementwise operations: add even elements with odd elements
- Time them against their pure python counterparts using `%timeit`.
- Generate:
 - `[2**0, 2**1, 2**2, 2**3, 2**4]`
 - `a_j = 2^(3*j) - j`

Other operations

Comparisons:

```
In [70]: 1 import numpy as np
```

```
In [71]: 1 a = np.array([1, 2, 3, 4])
2 b = np.array([4, 2, 2, 4])
3
4 print(a == b)
5 print(a > b)
```

```
[False True False True]
[False False True False]
```

Array-wise comparisons:

```
In [72]: 1 a = np.array([1, 2, 3, 4])
2 b = np.array([4, 2, 2, 4])
3 c = np.array([1, 2, 3, 4])
4
5 print(np.array_equal(a, b))
6 print(np.array_equal(a, c))
```

```
False
True
```

Logical operations:

```
In [73]: 1 a = np.array([1, 1, 1, 0], dtype=bool)
2 b = np.array([1, 0, 1, 0], dtype=bool)
3
4 print(np.logical_or(a, b))
5 print(np.logical_and(a, b))
```

```
[ True  True  True False]
[ True False  True False]
```

Transcendental functions:

```
In [74]: 1 a = np.arange(1,5)
2
3 print(np.sin(a))
4 print(np.log(a))
5 print(np.exp(a))
```

```
[ 0.84147098  0.90929743  0.14112001 -0.7568025 ]
[0.          0.69314718  1.09861229  1.38629436]
[ 2.71828183  7.3890561  20.08553692 54.59815003]
```

Shape mismatches

```
In [75]: 1 a = np.arange(4)
          2 print(a.shape)
          3 a + np.array([1, 2]) # Gives an error!

(4,)
```

```
ValueError                                Traceback (most recent call last)
~\AppData\Local\Temp\ipykernel_16448\2497161380.py in <module>
      1 a = np.arange(4)
      2 print(a.shape)
----> 3 a + np.array([1, 2]) # Gives an error!

ValueError: operands could not be broadcast together with shapes (4,) (2,)
```

Transposition:

```
In [76]: 1 a = np.triu(np.ones((3, 3)),2) # see help(np.triu)
          2 a = np.tril(np.ones((3, 3)), -2)
          3 print(a)
          4 print()
          5 print(a.T)

[[0. 0. 0.]
 [0. 0. 0.]
 [1. 0. 0.]]

[[0. 0. 1.]
 [0. 0. 0.]
 [0. 0. 0.]]
```

Note: The transposition is a view

As a result, the following code is **wrong**, and will **not make a matrix symmetric**:

```
In [77]: 1 a += a.T
          2 print(a)

[[0. 0. 1.]
 [0. 0. 0.]
 [1. 0. 0.]]
```

It will work for small arrays (because of buffering) but fail for large one, in unpredictable ways.

Basic reductions

Computing sums

```
In [78]: 1 x = np.array([1, 2, 3, 4])
          2
          3 print(np.sum(x))
          4 print(x.sum())
          5 print(x.cumsum())

10
10
[ 1  3  6 10]
```

Sum by rows and by columns:

```
In [79]: 1 x = np.array([[1, 1], [2, 2]])
          2 x
```

```
Out[79]: array([[1, 1],
                [2, 2]])
```

```
In [80]: 1 x.sum(axis=0) # columns (first dimension)
```

```
Out[80]: array([3, 3])
```

```
In [81]: 1 x[:, 0].sum(), x[:, 1].sum() # alternative manual approach
```

```
Out[81]: (3, 3)
```

```
In [82]: 1 x.sum(axis=1) # rows (second dimension)
```

```
Out[82]: array([2, 4])
```

```
In [83]: 1 x[0, :].sum(), x[1, :].sum() # alternative manual approach
```

```
Out[83]: (2, 4)
```

Same idea in higher dimensions:

```
In [84]: 1 x = np.random.rand(2, 2, 2)
2
3 print(x)
4 print(x.sum(axis=2)[0, 1])
```

```
[[[0.21233911 0.18182497]
  [0.18340451 0.30424224]]
```

```
 [[0.52475643 0.43194502]
  [0.29122914 0.61185289]]]
0.48764675281297154
```

```
In [85]: 1 print(x[0, 1, :].sum()) # alternative approach
```

```
0.48764675281297154
```

Other reductions

Extrema:

```
In [86]: 1 x = np.array([1, 3, 2])
2
3 print(x.min()) # minimum value
4 print(x.max()) # maximum value
5 print()
6 print(x.argmin()) # index of minimum value
7 print(x.argmax()) # index of maximum value
```

```
1
3

0
1
```

Logical operations:

```
In [87]: 1 print(np.all([True, True, False]))
2
3 print(np.any([True, True, False]))
```

```
False
True
```

Note: Can be used for array comparisons:

```
In [88]: 1 a = np.zeros((100, 100))
2
3 print(np.any(a != 0))
4 print(np.all(a == 0))
5
6 a = np.array([1, 2, 3, 2])
7 b = np.array([2, 2, 3, 2])
8 c = np.array([6, 4, 4, 5])
9 print((a <= b) & (b <= c)).all()
```

```
False
True
True
```

Statistics:

```
In [89]: 1 x = np.array([1, 2, 3, 1])
2 y = np.array([[1, 2, 3], [5, 6, 1]])
3
4 print(x.mean())
5 print(np.median(x))
6
7 print(y)
8 print(np.median(y,axis=-1)) # last axis
9
10 print(x.std())           # full population standard dev.
11 print(x.cumsum())

1.75
1.5
[[1 2 3]
 [5 6 1]]
[2. 5.]
0.82915619758885
[1 3 6 7]
```

```
In [90]: 1 # np.median? # interactive help
```

Exercise: Reductions

- What is the difference between `sum` and `cumsum` ?

Worked example: Data statistics

Data in `populations.txt` describes the populations of hares and lynxes (and carrots) in northern Canada during 20 years.

First view the data:

```
In [91]: 1 !cat data/populations.txt

# year hare lynx carrot
1900 30e3 4e3 48300
1901 47.2e3 6.1e3 48200
1902 70.2e3 9.8e3 41500
1903 77.4e3 35.2e3 38200
1904 36.3e3 59.4e3 40600
1905 20.6e3 41.7e3 39800
1906 18.1e3 19e3 38600
1907 21.4e3 13e3 42300
1908 22e3 8.3e3 44500
1909 25.4e3 9.1e3 42100
1910 27.1e3 7.4e3 46000
1911 40.3e3 8e3 46800
1912 57e3 12.3e3 43800
1913 76.6e3 19.5e3 40900
1914 52.3e3 45.7e3 39400
1915 19.5e3 51.1e3 39000
1916 11.2e3 29.7e3 36700
1917 7.6e3 15.8e3 41800
1918 14.6e3 9.7e3 43300
1919 16.2e3 10.1e3 41300
1920 24.7e3 8.6e3 47300
```

Now, load the data into a NumPy array:


```
In [92]: 1 data = np.loadtxt('data/populations.txt')
2
3 year, hares, lynxes, carrots = data.T # trick: columns to variables
4
5 year

Out[92]: array([1900., 1901., 1902., 1903., 1904., 1905., 1906., 1907., 1908.,
1909., 1910., 1911., 1912., 1913., 1914., 1915., 1916., 1917.,
1918., 1919., 1920.])
```

Then, plot it:

```
In [93]: 1 from matplotlib import pyplot as plt
2 plt.axes([0.2, 0.1, 0.5, 0.8])
3 plt.plot(year, hares, year, lynxes, year, carrots)
4 plt.legend(('Hare', 'Lynx', 'Carrot'), loc=(1.05, 0.5))
```

Out[93]: <matplotlib.legend.Legend at 0x1e446fb5580>



The mean populations over time:

```
In [94]: 1 populations = data[:, 1:]
2 populations.mean(axis=0)
```

Out[94]: array([34080.95238095, 20166.66666667, 42400.])

The sample standard deviations:

```
In [95]: 1 populations.std(axis=0)
```

Out[95]: array([20897.90645809, 16254.59153691, 3322.50622558])

Which species has the highest population each year?:

```
In [96]: 1 np.argmax(populations, axis=1)
```

Out[96]: array([2, 2, 0, 0, 1, 1, 2, 2, 2, 2, 2, 2, 0, 0, 0, 1, 2, 2, 2, 2, 2],
dtype=int64)

Array shape manipulation

Flattening

```
In [97]: 1 a = np.array([[1, 2, 3], [4, 5, 6]])
2 print(a)
3 print(a.ravel())
4 print("\n")
5 print(a)
6 b = a.T
7 print(b)
8 print(b.ravel())
```

```
[[1 2 3]
 [4 5 6]]
[1 2 3 4 5 6]
```

```
[[1 2 3]
 [4 5 6]]
[[1 4]
 [2 5]
 [3 6]]
[1 4 2 5 3 6]
```

```
In [98]: 1 a.T.ravel()
```

Out[98]: array([1, 4, 2, 5, 3, 6])

Note: In higher dimensions: last dimensions ravel out "first".

Reshaping

The inverse operation to flattening:

```
In [99]: 1 print(a)
```

```
[[1 2 3]
 [4 5 6]]
```

```
In [100]: 1 print(a.ravel())
```

```
[1 2 3 4 5 6]
```

```
In [101]: 1 a.reshape(3,2)
          2 print(a)
```

```
[[1 2 3]
 [4 5 6]]
```

```
In [102]: 1 print(a.shape)
```

```
(2, 3)
```

```
In [103]: 1 b = a.ravel()
          2 print(b)
          3 print(b.shape)
          4 b = b.reshape((3, 2))
          5 print(b)
```

```
[1 2 3 4 5 6]
(6,)
[[1 2]
 [3 4]
 [5 6]]
```

```
In [104]: 1 a.reshape((3, -1))  # unspecified (-1) value is inferred
```

```
Out[104]: array([[1, 2],
                 [3, 4],
                 [5, 6]])
```

Note: `ndarray.reshape` may return a view (cf `help(np.reshape)`), or copy

Beware: `reshape` may also return a copy!

```
In [105]: 1 a = np.zeros((3, 2))
          2 b = a.T.reshape(3*2)
          3 b[0] = 9
          4 print(a)
```

```
[[0. 0.]
 [0. 0.]
 [0. 0.]]
```

To understand this you need to learn more about the memory layout of a numpy array.

Adding a dimension

Indexing with the `np.newaxis` object allows us to add an axis to an array (you have seen this already above in the broadcasting section):

```
In [106]: 1 import numpy as np
          2 z = np.array([1, 2, 3])
          3 print(z)
          4 z.shape
```

```
[1 2 3]
```

```
Out[106]: (3,)
```



```
In [107]: 1 a = z[:, np.newaxis]
          2 print(a)
          3 print(a.shape)
```

```
[[1]
 [2]
 [3]]
(3, 1)
```

```
In [108]: 1 m = z[np.newaxis, :]
          2 print(m)
```

```
[[1 2 3]]
```

```
In [109]: 1 print(m.shape)
```

```
(1, 3)
```

Dimension shuffling

```
In [110]: 1 import numpy as np
```

```
In [111]: 1 a = np.arange(4*3*2).reshape(4, 3, 2)
          2 print(a.shape)
          3 print(a)
```

```
(4, 3, 2)
[[[ 0  1]
 [ 2  3]
 [ 4  5]]
```

```
 [[ 6  7]
 [ 8  9]
 [10 11]]
```

```
[[12 13]
 [14 15]
 [16 17]]
```

```
[[18 19]
 [20 21]
 [22 23]]]
```

```
In [112]: 1 print(a[2, 2, 1])
```

```
17
```

```
In [113]: 1 b = a.transpose(1, 2, 0)
          2 print(b.shape)
          3 print(b)
```

```
(3, 2, 4)
[[[ 0  6 12 18]
 [ 1  7 13 19]]
```

```
 [[ 2  8 14 20]
 [ 3  9 15 21]]
```

```
[[ 4 10 16 22]
 [ 5 11 17 23]]]
```

```
In [114]: 1 c = a.transpose(1,0,2)
          2 print(c.shape)
          3 print(c)
```

```
(3, 4, 2)
[[[ 0  1]
   [ 6  7]
   [12 13]
   [18 19]]
```

```
[[ 2  3]
 [ 8  9]
 [14 15]
 [20 21]]
```

```
[[ 4  5]
 [10 11]
 [16 17]
 [22 23]]]
```

```
In [115]: 1 print(b[2, 1, 0])
```

```
5
```

Also creates a view:

```
In [116]: 1 b[2, 1, 0] = -1
          2 print(a[0, 2, 1])
```

```
-1
```

Resizing

Size of an array can be changed with `ndarray.resize` :

```
In [117]: 1 a = np.arange(4)
          2 print(a)
          3 a.resize((8,))
          4 print(a)
```

```
[0 1 2 3]
[0 1 2 3 0 0 0 0]
```

However, it must not be referred to somewhere else:

```
In [118]: 1 b = a
          2 a.resize((4,)) # Gives an error!
```

```
-----
ValueError                                Traceback (most recent call last)
~\AppData\Local\Temp\ipykernel_16448\2941388788.py in <module>
      1 b = a
----> 2 a.resize((4,)) # Gives an error!
```

```
ValueError: cannot resize an array that references or is referenced
by another array in this way.
Use the np.resize function or refcheck=False
```

Exercise: Shape manipulations

- Look at the docstring for `reshape`, especially the notes section which has some more information about copies and views.
- Use `flatten` as an alternative to `ravel`. What is the difference? (Hint: check which one returns a view and which a copy)
- Experiment with `transpose` for dimension shuffling.

Sorting data

Sorting along an axis:

```
In [119]: 1 a = np.array([[4, 3, 5], [1, 2, 1]])
          2 print(a)
          3 b = np.sort(a, axis=-1)
          4 print(b)
```

```
[[4 3 5]
 [1 2 1]]
[[3 4 5]
 [1 1 2]]
```

```
In [120]: 1 # np.sort? # Interactive help!
```

Note: Sorts each row separately!

In-place sort:

```
In [121]: 1 a.sort(axis=1)
          2 print(a)
```

```
[[3 4 5]
 [1 1 2]]
```

Finding minima and maxima:

```
In [122]: 1 a = np.array([4, 3, 1, 2])
          2 j_max = np.argmax(a)
          3 j_min = np.argmin(a)
          4 print(j_max, j_min)
```

```
0 2
```

Exercise: Sorting

- Try both in-place and out-of-place sorting.
- Try creating arrays with different dtypes and sorting them.
- Use `all` or `array_equal` to check the results.
- Look at `np.random.shuffle` for a way to create sortable input quicker.
- Combine `ravel`, `sort` and `reshape`.
- Look at the `axis` keyword for `sort` and rewrite the previous exercise.

Summary

What do you need to know to get started?

- Know how to create arrays : `array` , `arange` , `ones` , `zeros` .
- Know the shape of the array with `array.shape` , then use slicing to obtain different views of the array: `array[:2]` , etc. Adjust the shape of the array using `reshape` or flatten it with `ravel` .
- Obtain a subset of the elements of an array and/or modify their values with masks

```
In [123]: 1 print(a)
```

```
[4 3 1 2]
```

```
In [124]: 1 a[a < 2] = 5
          2 print(a)
```

```
[4 3 5 2]
```

- Know miscellaneous operations on arrays, such as finding the mean or max (`array.max()` , `array.mean()`). No need to retain everything, but have the reflex to search in the documentation (online docs, `help()` , `lookfor()`)!!
- For advanced use: master the indexing with arrays of integers, as well as broadcasting. Know more NumPy functions to handle various array operations.

More elaborate arrays

More data types

Casting

“Bigger” type wins in mixed-type operations:

```
In [125]: 1 np.array([1, 2, 3]) + 1.0
```

```
Out[125]: array([2., 3., 4.])
```

Assignment never changes the type!

```
In [126]: 1 a = np.array([1, 2, 3])
          2 print(a.dtype)
          3 print(a)
```

```
int32
[1 2 3]
```

```
In [127]: 1 a[0] = 1.9      # <-- float is truncated to integer
          2 print(a)
```

```
[1 2 3]
```

Forced casts:

```
In [128]: 1 a = np.array([1.7, 1.2, 1.6])
          2 b = a.astype(int) # <-- truncates to integer
          3 print(b)
```

```
[1 1 1]
```

Rounding:

```
In [129]: 1 a = np.array([1.2, 1.5, 1.6, 2.5, 3.5, 4.5])
          2 b = np.around(a)
          3 print(b)           # still floating-point
```

```
[1.  2.  2.  2.  4.  4.]
```

```
In [130]: 1 c = np.around(a).astype(int)
          2 print(c)
```

```
[1 2 2 2 4 4]
```

In this session, we learnt some of the key features of NumPy. However, numPy is a very large library, and has countless other features. You can find more about it in the link below!

http://www.scipy-lectures.org/intro/numpy/elaborate_arrays.html (http://www.scipy-lectures.org/intro/numpy/elaborate_arrays.html).